

A Simple LJ Many-Body Simulator Optimization and Parallelization


Dr. Axel Kohlmeyer

Associate Dean for High-Performance Computing
Associate Director, ICMS
College of Science and Technology
Temple University, Philadelphia

<http://sites.google.com/site/akohlmey/>

a.kohlmeyer@temple.edu

The LJ Model for Liquid Argon



- Cubic box of particles with a Lennard-Jones type pairwise additive interaction potential

$$V = \sum_{i,j} \begin{cases} 4\epsilon \left[\left(\frac{\sigma}{r_{ij}} \right)^{12} - \left(\frac{\sigma}{r_{ij}} \right)^6 \right], & r_{ij} < r_c \\ 0, & r_{ij} \geq r_c \end{cases}$$

- Periodic boundary conditions to avoid surface effects

What Do We Need to Program?

1. Read in parameters and initial status and compute what is missing (e.g. accelerations)
2. Integrate Equations of motion with Velocity Verlet for a given number of steps
 - a) Propagate all velocities for half a step
 - b) Propagate all positions for a full step
 - c) Compute forces on all atoms to get accelerations
 - d) Propagate all velocities for half a step
 - e) Output intermediate results, if needed

Initial Serial Code: Velocity Verlet

```
void velverlet(mdsys_t *sys) {  
    for (int i=0; i<sys->natoms; ++i) {  
        sys->vx[i] += 0.5*sys->dt / mvsq2e * sys->fx[i] / sys->mass;  
        sys->vy[i] += 0.5*sys->dt / mvsq2e * sys->fy[i] / sys->mass;  
        sys->vz[i] += 0.5*sys->dt / mvsq2e * sys->fz[i] / sys->mass;  
        sys->rx[i] += sys->dt*sys->vx[i];  
        sys->ry[i] += sys->dt*sys->vy[i];  
        sys->rz[i] += sys->dt*sys->vz[i];  
    }  
}
```

```
    force(sys);
```

```
    for (int i=0; i<sys->natoms; ++i) {  
        sys->vx[i] += 0.5*sys->dt / mvsq2e * sys->fx[i] / sys->mass;  
        sys->vy[i] += 0.5*sys->dt / mvsq2e * sys->fy[i] / sys->mass;  
        sys->vz[i] += 0.5*sys->dt / mvsq2e * sys->fz[i] / sys->mass;  
    }  
}
```


Initial Code: Force Calculation

```
for(i=0; i < (sys->natoms); ++i) {  
    for(j=0; j < (sys->natoms); ++j) {  
        if (i==j) continue;
```

```
        rx=pbcc(sys->rx[i] - sys->rx[j], 0.5*sys->box);  
        ry=pbcc(sys->ry[i] - sys->ry[j], 0.5*sys->box);  
        rz=pbcc(sys->rz[i] - sys->rz[j], 0.5*sys->box);  
        r = sqrt(rx*rx + ry*ry + rz*rz);
```

Compute distance
between atoms i & j

```
        if (r < sys->rcut) {
```

Compute energy and force

```
            ffac = -4.0*sys->epsilon*(-12.0*pow(sys->sigma/r,12.0)/r  
                +6*pow(sys->sigma/r,6.0)/r);  
            sys->epot += 0.5*4.0*sys->epsilon*(pow(sys->sigma/r,12.0)  
                -pow(sys->sigma/r,6.0));
```

```
            sys->fx[i] += rx/r*ffac;  
            sys->fy[i] += ry/r*ffac;  
            sys->fz[i] += rz/r*ffac;
```

Add force contribution
of atom j on atom i

```
        }  
    }
```

How Well Does it Work?

- Compiled with:

gcc -o ljmd.x ljmd.c -lm

Test input: 108 atoms, 10000 steps: 49s

Let us get a profile:

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
73.70	13.87	13.87	10001	1.39	1.86	force
24.97	18.57	4.70	346714668	0.00	0.00	pbc
0.96	18.75	0.18				main
0.37	18.82	0.07	10001	0.01	0.01	ekin
0.00	18.82	0.00	30006	0.00	0.00	azzero
0.00	18.82	0.00	101	0.00	0.00	output
0.00	18.82	0.00	12	0.00	0.00	getline

Compiler Optimization

- Use of pbc() is convenient, but costs 25%
=> compiling with -O3 should inline it
- Loops should be unrolled for superscalar CPUs
=> compiling with -O2 or -O3 should do it for us

Time now: 39s (1.3x faster) *Only a bit faster*

- Now try some more optimization options:
-ffast-math -fexpensive-optimizations -msse3

Time now: 10s (4.9x faster) *Much better!*

- Compare to LAMMPS: 3.6s => need to do more

Now Modify the Code

- Use physics! Newton's 3rd law: $F_{ij} = -F_{ji}$

```
for(i=0; i < (sys->natoms)-1; ++i) {  
    for(j=i+1; j < (sys->natoms); ++j) {  
        rx=pbcc(sys->rx[i] - sys->rx[j], 0.5*sys->box);  
        ry=pbcc(sys->ry[i] - sys->ry[j], 0.5*sys->box);  
        rz=pbcc(sys->rz[i] - sys->rz[j], 0.5*sys->box);  
        r = sqrt(rx*rx + ry*ry + rz*rz);  
        if (r < sys->rcut) {  
            ffac = -4.0*sys->epsilon*(-12.0*pow(sys->sigma/r,12.0)/r  
                                +6*pow(sys->sigma/r,6.0)/r);  
            sys->epot += 0.5*4.0*sys->epsilon*(pow(sys->sigma/r,12.0)  
                                -pow(sys->sigma/r,6.0));  
            sys->fx[i] += rx/r*ffac;        sys->fx[j] -= rx/r*ffac;  
            sys->fy[i] += ry/r*ffac;        sys->fy[j] -= ry/r*ffac;  
            sys->fz[i] += rz/r*ffac;        sys->fz[j] -= rz/r*ffac;  
        }  
    }  
}
```

Time now: 5.4s (9.0x faster) Another big improvement

More Modifications

- Avoid expensive math: pow(), sqrt(), division

```
c12=4.0*sys->epsilon*pow(sys->sigma,12.0);
c6 =4.0*sys->epsilon*pow(sys->sigma, 6.0);
rcsq = sys->rcut * sys->rcut;
for(i=0; i < (sys->natoms)-1; ++i) {
    for(j=i+1; j < (sys->natoms); ++j) {
        rx=pbcc(sys->rx[i] - sys->rx[j], 0.5*sys->box);
        ry=pbcc(sys->ry[i] - sys->ry[j], 0.5*sys->box);
        rz=pbcc(sys->rz[i] - sys->rz[j], 0.5*sys->box);
        rsq = rx*rx + ry*ry + rz*rz;
        if (rsq < rcsq) {
            double r6,rinv; rinv=1.0/rsq; r6=rinv*rinv*rinv;
            ffac = (12.0*c12*r6 - 6.0*c6)*r6*rinv;
            sys->epot += r6*(c12*r6 - c6);
            sys->fx[i] += rx*ffac; sys->fx[j] -= rx*ffac;
            sys->fy[i] += ry*ffac; sys->fy[j] -= ry*ffac;
            sys->fz[i] += rz*ffac; sys->fz[j] -= rz*ffac;
        }
    }
}
```

=> 108 atoms: 4.0s (12.2x faster) **still worth it**

What if optimization is not enough?

- Having linear scaling is nice, but twice the system size is still twice the work
=> Parallelization
- Simple MPI parallelization first
 - MPI is “share nothing” (replicated or distributed data)
 - Run the same code path with the same data but insert a few MPI calls
 - Broadcast positions from rank 0 to all before force()
 - Compute forces on different atoms for each rank
 - Collect (reduce) forces from all to rank 0 after force()

Replicated Data MPI Version

```
static void force(mdsys_t *sys) {  
    double epot=0.0;  
    azzero(sys->cx,sys->natoms); azzero(sys->cy,sys->natoms); azzero(sys->cz,sys->natoms);  
    MPI_Bcast(sys->rx, sys->natoms, MPI_DOUBLE, 0, sys->mpicomm);  
    MPI_Bcast(sys->ry, sys->natoms, MPI_DOUBLE, 0, sys->mpicomm);  
    MPI_Bcast(sys->rz, sys->natoms, MPI_DOUBLE, 0, sys->mpicomm);  
    for (i=0; i < sys->natoms-1; i += sys->nsz) {  
        ii = i + sys->mpirank;  
        if (ii >= (sys->natoms - 1)) break;  
        for (j=i+1; j < sys->natoms; ++j) {  
            [...]   
                sys->cy[j] -= ry*ffac;  
                sys->cz[j] -= rz*ffac;  
        } }  
        MPI_Reduce(sys->cx, sys->fx, sys->natoms, MPI_DOUBLE, MPI_SUM, 0, sys->mpicomm);  
        MPI_Reduce(sys->cy, sys->fy, sys->natoms, MPI_DOUBLE, MPI_SUM, 0, sys->mpicomm);  
        MPI_Reduce(sys->cz, sys->fz, sys->natoms, MPI_DOUBLE, MPI_SUM, 0, sys->mpicomm);  
        MPI_Reduce(&epot, &sys->epot, 1, MPI_DOUBLE, MPI_SUM, 0, sys->mpicomm);  
    }  
}
```

- Easy to implement, but lots of communication

OpenMP Parallelization

- OpenMP is directive based
=> code (can) work without them
- OpenMP can be added incrementally
- OpenMP only works in shared memory
=> multi-core processors
- OpenMP hides the calls to a threads library
=> less flexible, but less programming
- **Caution:** write access to shared data can easily lead to race conditions

Naive OpenMP Version

```
#if defined(_OPENMP)
#pragma omp parallel for default(shared) \
    private(i) reduction(+:epot)
#endif
for(i=0; i < (sys->natoms)-1; ++i) {
    double rx1=sys->rx[i];
    double ry1=sys->ry[i];
    double rz1=sys->rz[i];
    [...]

    {
        sys->fx[i] += rx*ffac;
        sys->fy[i] += ry*ffac;
        sys->fz[i] += rz*ffac;
        sys->fx[j] -= rx*ffac;
        sys->fy[j] -= ry*ffac;
        sys->fz[j] -= rz*ffac;
    }
}
```

Each thread will
work on different
values of “i”

Race condition: “i” will be unique for each thread, but not “j”
=> multiple threads may write to the same location concurrently

Naive OpenMP Version

```
#if defined(_OPENMP)
#pragma omp parallel for default(shared) \
private(i) reduction(+:epot)
#endif
```

```
for(i=0; i < (sys->natoms)-1; ++i) {
    double rx1=sys->rx[i];
    double ry1=sys->ry[i];
    double rz1=sys->rz[i];
    [...]
```

```
#if defined(_OPENMP)
#pragma omp critical
#endif
```

Each thread will
work on different
values of “i”

The “critical” directive will let only
one thread at a time execute this block

Timings (108 atoms): {

1 thread: 4.2s

2 threads: 7.1s

4 threads: 7.7s

8 threads: 8.6s

```
sys->fx[i] += rx*ffac;
sys->fy[i] += ry*ffac;
sys->fz[i] += rz*ffac;
sys->fx[j] -= rx*ffac;
sys->fy[j] -= ry*ffac;
sys->fz[j] -= rz*ffac;
```

This is making it
slower not faster!

} LJMD Simulation Code

OpenMP Improvements

- Use **omp atomic** to protect one instruction
=> faster, but requires hardware support
108: 1T: 6.3s, 2T: 5.0s, 4T: 4.4s, 8T: 4.2s
2916: 1T: 126s, 2T: 73s, 4T: 48s, 8T: 26s
=> some speedup, but noticable overhead
=> serial is faster than OpenMP with 1T
- Don't use Newton's 3rd Law => no race condition
108: 1T: 6.5s, 2T: 3.7s, 4T: 2.3s, 8T: 2.1s
2916: 1T: 213s, 2T: 106s, 4T: 53s, 8T: 21s
=> better scaling, but we lose 2x serial speed

MPI-like Approach with OpenMP

```
#if defined(_OPENMP)
#pragma omp parallel reduction(+:epot)
#endif
```

```
    { double *fx, *fy, *fz;
```

```
#if defined(_OPENMP)
```

```
    int tid=omp_get_thread_num();
```

```
#else
```

```
    int tid=0;
```

```
#endif
```

```
    fx=sys->fx + (tid*sys->natoms); azzero(fx,sys->natoms);
```

```
    fy=sys->fy + (tid*sys->natoms); azzero(fy,sys->natoms);
```

```
    fz=sys->fz + (tid*sys->natoms); azzero(fz,sys->natoms);
```

```
    for(int i=0; i < (sys->natoms -1); i += sys->nthreads) {
```

```
        int ii = i + tid;
```

```
        if (ii >= (sys->natoms -1)) break;
```

```
        rx1=sys->rx[ii];
```

```
        ry1=sys->ry[ii];
```

```
        rz1=sys->rz[ii];
```

Thread Id is like MPI rank

sys->fx holds storage for one full fx array for each thread => race condition is eliminated.

MPI-like Approach with OpenMP (2)

- We need to write our own reduction:

```
#if defined (_OPENMP)
#pragma omp barrier
#endif
```

Need to make certain, all threads
are done with computing forces

```
i = 1 + (sys->natoms / sys->nthreads);
fromidx = tid * i;
toidx = fromidx + i;
if (toidx > sys->natoms) toidx = sys->natoms;
```

```
for (i=1; i < sys->nthreads; ++i) {
    int offs = i*sys->natoms;
    for (int j=fromidx; j < toidx; ++j) {
        sys->fx[j] += sys->fx[offs+j];
        sys->fy[j] += sys->fy[offs+j];
        sys->fz[j] += sys->fz[offs+j];
    }
}
```

Use threads to
parallelize the
reductions

More OpenMP Timings

- The **omp parallel** region timings
108: 1T: 3.5s, 2T: 2.5s, 4T: 2.2s, 8T: 2.5s
2916: 1T: 103s, 2T: 53s, 4T: 19s, 8T: 10s
=> better speedup, 1T is about as fast as serial
=> scaling like no 3rd law, but speed 2x as fast
- This approach also works with cell lists:
108: 1T: 4.3s, 2T: 3.1s, 4T: 2.4s, 8T: 2.9s
2916: 1T: 28s, 2T: 15s, 4T: 8.9s, 8T: 4.1s
=> 6.8x speedup with 8 threads.
62x faster than original code with 2916 atoms

Hybrid OpenMP/MPI Version

- With multi-core nodes, communication between MPI tasks becomes a problem
 - => all communication has to use one link
 - => reduced bandwidth, increased latency
- OpenMP and MPI parallelization are orthogonal and can be used at the same time
 - Caution:** don't call MPI from threaded region
- Parallel region OpenMP version is very similar to MPI version, so that would be easy to merge

Hybrid OpenMP/MPI Kernel

- MPI tasks are like GPU thread blocks
- Need to reduce forces/energies first across threads and then across all MPI tasks

[...]

```
incr = sys->mpisize * sys->nthreads;  
/* self interaction of atoms in cell */  
for(n=0; n < sys->ncell; n += incr) {  
    int i,j;  
    const cell_t *c1;  
  
    i = n + sys->mpirank*sys->nthreads + tid;  
    if (i >= sys->ncell) break;  
    c1=sys->clist + i;  
  
    for (j=0; j < c1->natoms-1; ++j) {
```

[...]

Hybrid OpenMP/MPI Timings

2916 atoms system:

78732 atoms system:

Cell list serial code: 18s

50.1s

16 MPI x 1 Threads: 14s

19.8s

8 MPI x 2 Threads: 5.5s

8.9s

4 MPI x 4 Threads: 4.3s

8.2s

2 MPI x 8 Threads: 4.0s

7.3s

=> Best speedup: 4.5x

6.9x

=> Total speedup: **185x**

333x

- Replicated data MPI is simple to implement but does not parallelize well for this kind of code

Two nodes with 2x quad-core

A Simple LJ Many-Body Simulator Optimization and Parallelization

Dr. Axel Kohlmeyer

Associate Dean for High-Performance Computing
Associate Director, ICMS
College of Science and Technology
Temple University, Philadelphia

<http://sites.google.com/site/akohlmey/>

a.kohlmeyer@temple.edu